Track C facilitator run sheet

Coordinator reference | November 15, 2026 | Same version for all three cohorts

Before each round

- Rounds: 2:15 pm and 3:00 pm. Start a 45-minute countdown. Every coordinator delivers every segment independently. Have round-two packets ready before round one.
- Give everyone the four-page packet. Pages 2-3 are the exercise, page 4 the exit card. Collect the whole packet, not personal emails, generated PDFs or screenshots.
- Prepare a charged phone, tested app, timer, pens and paper. Check actual consent, export and disconnect screens before the event. Do not promise unverified controls. Hide notifications.
- Keep the fictional teaching examples on page 3 of this sheet available on paper. They are authored examples, not captured model results.

00-02 Intro | packet page 1

- Say: "You will use AI on familiar work, judge the result, and leave with a method for next week."
- "Use a personal account only. Email connection is optional. You may follow a willing neighbor, who holds their phone and chooses what is safe to show. Paper participation also counts."
- "Nothing confidential goes into a tool. Nothing is sent, paid, deleted or publicly shared. Do not read private messages aloud."
- Pick a timekeeper per table. Quick experience check; explain unfamiliar terms as needed. Do not move people between tables.

02-07 Tool landscape | page 1

- Drafting/explaining: turn rough points into a reply or explain unfamiliar text. Give context and review the answer.
- Finding/comparing: locate or organize information. Check source and date; a confident summary is not evidence itself.
- Acting across apps: agents may use connected tools or carry out steps. Ask what they can read/change and where approval happens. Chat assistants can also have action tools; product labels do not define access.
- Muse is the worked example; familiar chat tools can support text exercises. No feature, price or account entitlement claims from memory.
- Relatable example: rough notes into decisions and actions. An invented deadline can be beautifully written and still wrong. Choose the simplest tool and smallest access for the job.
- Desktop contrasts: brief mention only. No setup, account tour or new collaboration product in the core path.

07-12 Prompting | page 2

- Point to Context, Intent, Format, Restrictions. "Situation, job, shape, boundaries. Then define what good looks like."
- Weak: "Write an email about the review." Strong: "Our review moves from November 19 at 10 am to November 24 at 10 am. Ask for agenda confirmation by November 20. Write under 120 words; lead with the new date; use only these facts."
- Quality bar: correct date/time, one request/deadline, no invented reason or promise, concise. Fictional dates and quantities are fine; confidential figures are not.
- Circle one habit on page 2. Do not teach six separate mini-lessons. One relevant example is enough today.
- "Ask the model to check the draft, then check it yourself against the source. Its self-check is not verification."
- Before connection: "Read actual consent. Access may cover more than three selected messages. Decline if unsure. Personal inboxes can contain other people's sensitive data."

The 25 minute exercise

Packet pages 2-3 | Five parts, five minutes each | Owner always holds the phone

12-17 Part 1 Prompt it | page 2

- Choose a familiar task: email, summary, meeting-prep note or checklist. Keep sensitive text out of the app and the worksheet being collected.
- Write four parts and a checkable quality bar. Example: "Three actions with dates from my notes; flag missing dates rather than invent them."
- For inbox participants, choose an appropriate email task. If sensitive, change tasks before showing anything.
- Circulate: "Who is it for? What makes it useful? What must it not invent?" At minute 16, finish one usable prompt.

17-22 Part 2 Brief and draft | page 2

- Participants choosing inbox access select three non-sensitive messages. Ask: "For each, what do I need to know, when does it matter, and what do I do next? Use only the messages; flag missing dates."
- Check summary against source, then draft the most urgent suitable reply. No sending. Write sanitized Know / When / Next notes and key draft lines.
- Non-email participants run their chosen prompt. Those following a neighbor record their own observations. Owner controls what is visible; nobody browses another person's inbox.
- Setup limit: one minute. Then use a willing neighbor, appropriate text in an existing app, or paper. Do not tell people to accept every permission.

22-27 Part 3 Make a document | page 3

- Make one one-page reference: style guide, recurring-task checklist, role cheat sheet or "how to work with me" guide. Reuse the chosen task/material where possible.
- For the printed style-guide prompt, narrow "recent notes and emails" to selected non-sensitive samples. Do not invite unrestricted scanning. No suitable samples? Choose a checklist.
- Prompt: "Using only selected material, make a one-page [reference]. Good looks like: [three checks]. Separate facts from suggestions; flag missing information."
- Export to PDF if the tested path supports it. If export fails, use text on screen or a saved note. Do not spend the exercise fixing an exporter.

27-32 Part 4 Judge and improve | page 3

- 1 minute: partner suggests one prompt change aloud; owner enters it and creates version B of the same document.
- 2 minutes: show only safe output to a willing judge. Owner keeps the phone. A/B can stand for left/right. Judge checks the quality bar and records one reason.
- 1 minute: judge names the change that would make the weaker version useful. Look for invented facts, missing details, generic filler, needless length. Deleting weak output is valid.
- 1 minute: owner tunes once. If regeneration is slow, write the revision. One voluntary sanitized table example only if time remains.
- It is blinded only if the judge cannot identify the revision. Preference does not establish accuracy, model rankings or measured savings.

32-37 Part 5 Map the agent | page 3

- Fill Trigger -> AI step -> Human approval -> Destination for a repeatable chore. This is a sketch, not a live automation build.
- Example: selected new message -> summarize/draft -> I check facts, recipient and tone -> draft stays for approval. For a checklist task, map its weekly update instead.
- Add estimated value and an implication. Include review/fix time. Ask: "What may it read/change? Where must it stop?"
- Fast finishers add a stop rule: no sending, purchases or sharing without approval. At minute 37, turn to page 4.

Review and close

Packet page 4 | Finish at minute 45 | Keep safety and exit checks

37-42 Evaluation and guardrails

- 37-38: five checks. Does it do the job? Accurate against source? Fast including fixes? Appropriate access? Would you use it again for this task?
- 38-39: AI slop. "What is generic, invented, repetitive or unnecessary? What would you delete?" Judge against the written quality bar.
- 39-40: data sort. Public job posting: GREEN. Fictional meeting time: GREEN. Internal notes with unclear permission: YELLOW, check policy first. Client financial record: RED today. Password: RED. Ask for one reason.
- 40-41: permission scenario. A tool asks to read all mail and send messages to draft one reply. "What would you decline or narrow?" Read-only still exposes content. Draft generation does not itself need sending rights. Use only consent choices actually shown.
- 41-42: review connected access and show the tested disconnect path. If connection was only for this workshop, suggest revoking it now. Disconnecting does not delete already processed data. No one needs to show their account screen.
- Debrief: employer approval matters even on a paid account; summaries miss urgency; polished answers invent facts; the human owns what is used or sent.

42-45 Wrap-up and collection

- 42-43: write a next step and date. Measure minutes including fixes, factual errors or edits needed. Choose a task, not "use AI more."
- 43-44: Week 1 draft; Week 2 summarize/research; Week 3 workflow map. In week 4, repeat the best task and compare with the baseline. Use the next-step space for specifics.
- 44-45: check exit boxes against written work. Collect whole packets, not phones or digital outputs. Generated artifacts remain privately on participants' devices.
- Captains stack packets without reading personal entries aloud. Designate a coordinator for collection/storage. Do not upload participant writing to the public repo.

Paper example if access or network fails

- Fictional source: review moved from November 19, 10 am to November 24, 10 am. Agenda confirmation due November 20.
- Weak: "We are thrilled to announce an exciting evolution of our collaboration. The review is now November 25 due to supplier delays. We will send a gift to thank you."
- Better: "Our review is now November 24 at 10 am. Please confirm the agenda by November 20. Does the new time work for you?"
- Ask which meets the date/request/no-invention checks. Cross out the invented date, reason and gift. Write an improved prompt. Make a checklist: verify date, confirm agenda, review recipient, stop before send.
- These authored examples support paper practice; they are not evidence of a tested Muse run. Use them if personal screen sharing is unsuitable.

If time slips or something goes wrong

- Cut extra product examples, prompt rounds and table reports first. Keep permission guidance, source checks and exit step. One revision beats repeated exports.
- Unsafe material: stop kindly, do not repeat it, switch to appropriate content. Do not record the sensitive detail.
- Setup failure: willing neighbor with owner control, appropriate text in an existing tool, or paper. No payment, laptop setup or mandatory account connection.
- Round change: separate completed packets, set out fresh ones, reset timer and close private screens. Both rounds get the same version.
- Before calling this delivery ready: record permission/export/disconnect tests, timed rehearsal/fixes and each coordinator's approval. This sheet implies none of those checks are complete.